


# AI-Powered Predictive Trading Analytics Platform

David Alexander Colorado Rodríguez

Andres Felipe Martin Rodriguez

Databases II, Systems Engineering Curriculum Project, Universidad Distrital Francisco José de Caldas, Colombia, Bogotá D.C.



## 1. Introduction

The financial trading landscape has evolved dramatically with algorithmic trading now representing 70-80% of trading volume in major exchanges. This project addresses the challenges of modern trading systems by developing an AI-powered predictive trading analytics platform that integrates multiple trading strategies with real-time market analysis. Traditional database architectures struggle to simultaneously support high-frequency trading operations and complex analytical workloads, necessitating a novel approach for effective financial decision-making.

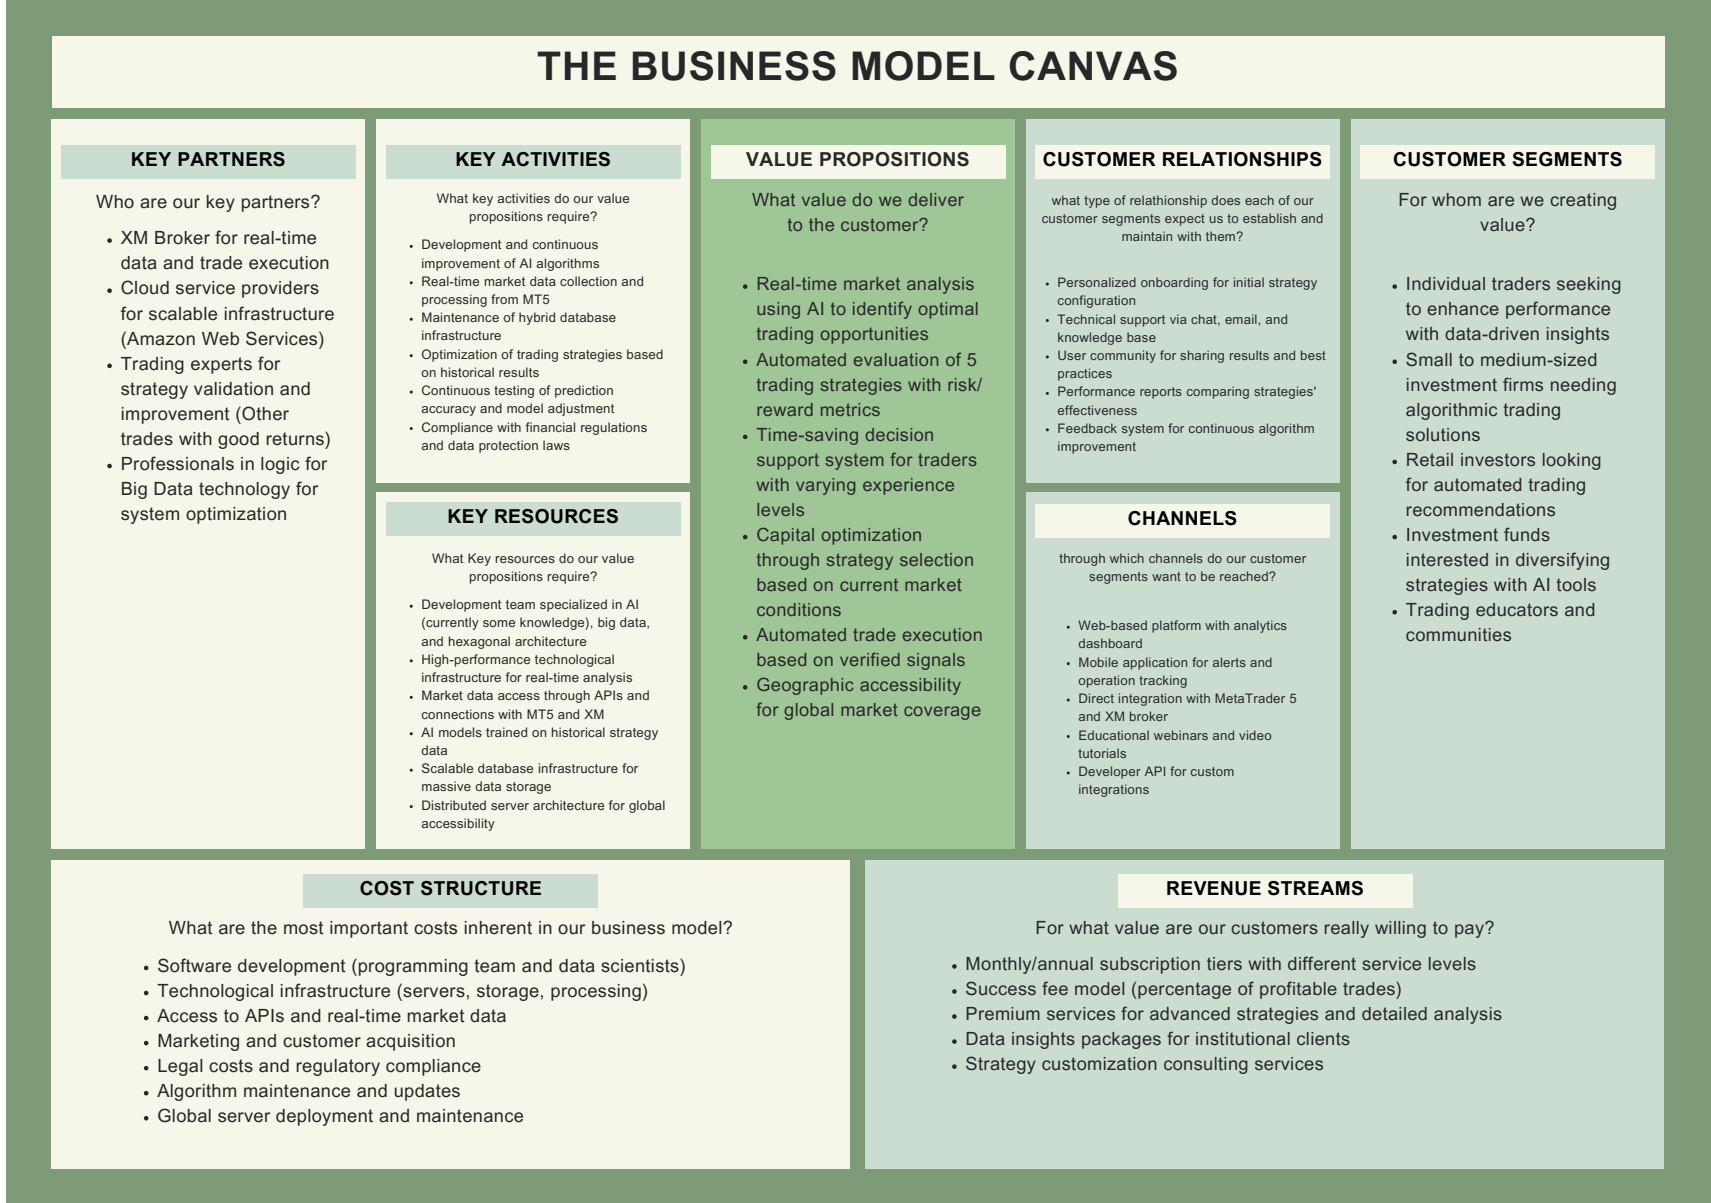
## 2. Goal

To design and implement a comprehensive trading analytics platform that leverages artificial intelligence to identify optimal trading opportunities across global markets while providing customizable strategy evaluation and automated trade execution. The system aims to deliver real-time insights for traders of varying experience levels while maintaining high performance, security, and regulatory compliance.

## 3. Proposed Solution

### Hybrid Database Architecture

- PostgreSQL:** Handles transactional operations requiring ACID compliance (user accounts, portfolios, trades)
- MongoDB:** Manages high-volume market data with time-series collections and compound indexes
- Snowflake:** Powers analytical workloads with separate compute warehouses for different query patterns

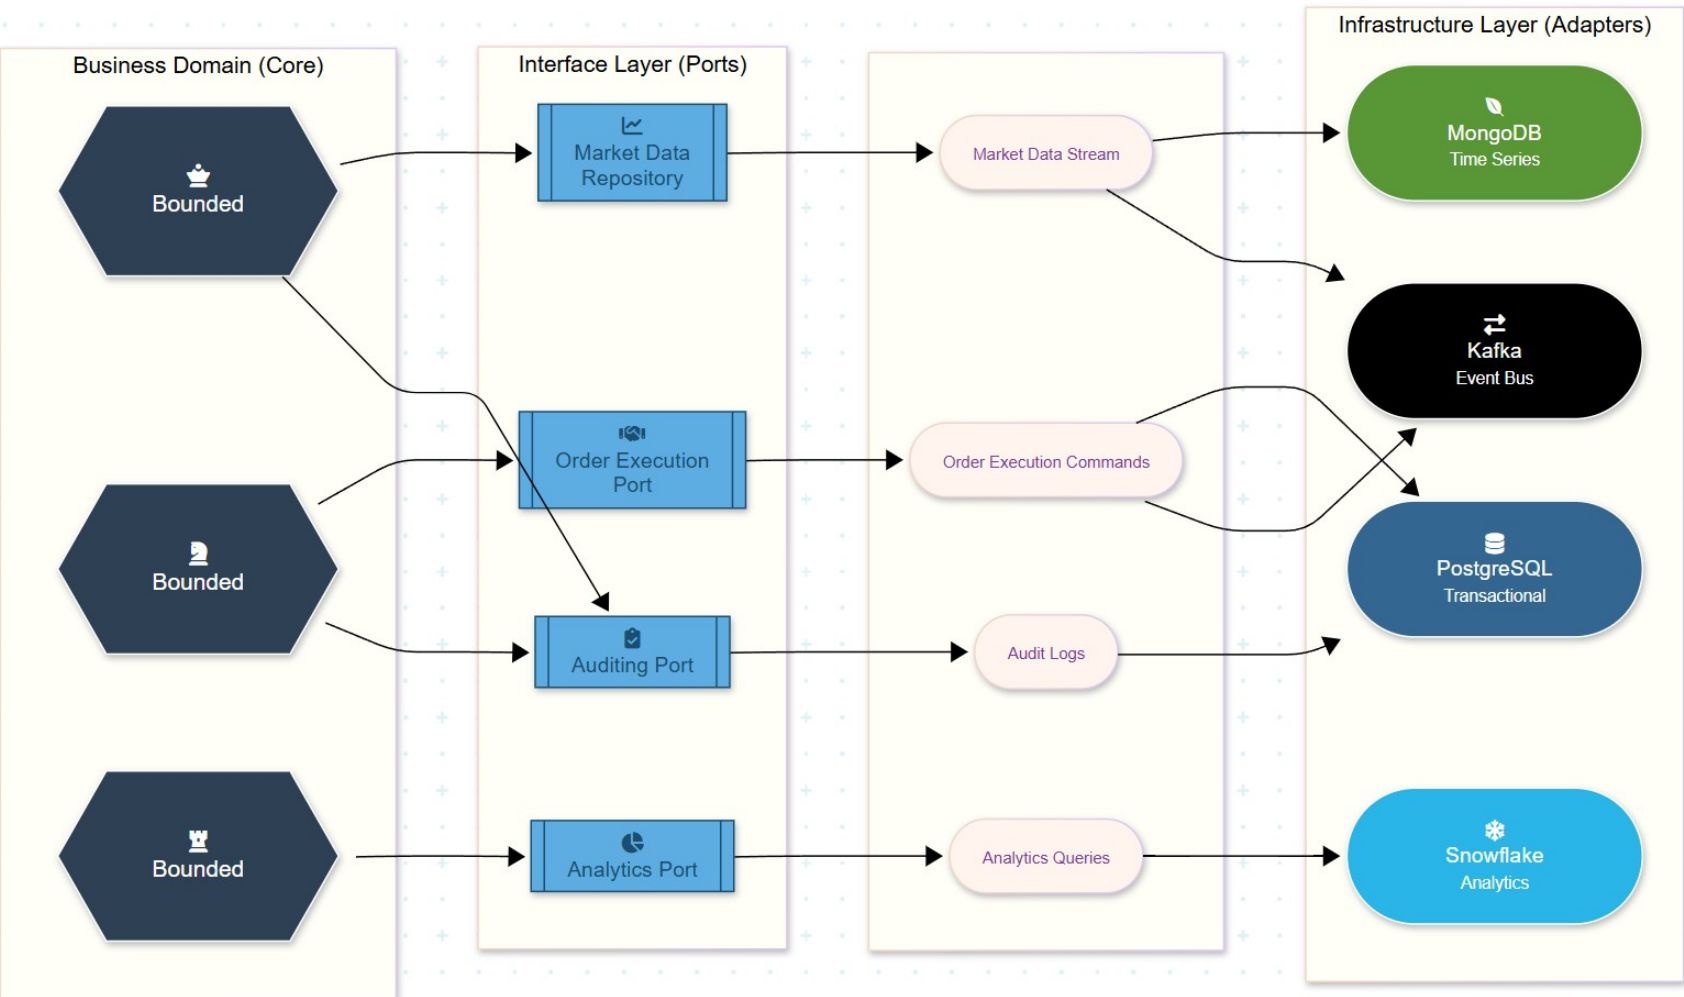


## Hexagonal Design Pattern

- Clear separation between domain logic and infrastructure components through well-defined interfaces
- Six primary domains: User Management, Portfolio Management, Strategy Configuration, Market Data, Trade Execution, and Performance Analytics

## Data Integration

- Apache Kafka serves as backbone for event-driven architecture
- Multi-stage market pipeline normalizes data from multiple providers (Yahoo Finance, Alpha Vantage)
- Geographic sharding ensures regulatory compliance across jurisdictions



## 4. Analysis

### Data Volume Projections

- 11,500 total users (10,000 individual traders, 500 investment firms, 1,000 retail investors)
- 130,000 daily alerts and 62,000 daily trades
- Monthly data volume: 32 GB (1.86 GB transactions + 30 GB market data)

### Security Implementation

- Row-level security policies in PostgreSQL restrict user access
- Field-level encryption for personally identifiable information
- Comprehensive audit logging system records all data access and modifications
- AES-256 encryption at rest and TLS 1.3 in transit

### Key Challenges Addressed

- Real-time data processing using message queuing and distributed processing
- Scalability through geographic replication and auto-scaling
- Regulatory compliance via region-specific data storage

## 5. Expected Outcomes

### Anticipated Performance Improvements

- Up to 37% reduction in alert latency:** Expected average latency around 412ms, improving over the 650ms industry benchmark
- Trade execution success rate potentially reaching 98.7%:** Surpassing the typical 95% industry standard
- Possible 33% decrease in storage costs:** Through optimization of approximately 1.2TB of compressed historical data

### Projected System Reliability

- Availability anticipated at 99.997% with no unplanned downtime
- Capability for graceful degradation under failure scenarios enabled by redundancy
- Disaster recovery expected to achieve PostgreSQL failover within 30 seconds

### Estimated Cost Efficiency

- Around 42% improvement anticipated via component-specific scaling and resource allocation
- Implementation of consumption-based pricing models for analytical workloads
- Automated tiering policies planned for managing aging market data

## 6. Future work

- Optimize synchronization mechanisms to reduce data staleness during peak loads
- Develop more efficient query planning for federated queries spanning PostgreSQL and Snowflake
- Explore federated learning for AI models to simplify regulatory compliance
- Enhance trading strategy automation based on historical performance metrics
- Implement more sophisticated market condition classification for improved strategy selection
- Develop advanced portfolio optimization models utilizing multi-factor risk analysis